



Houseloan-data-analysis-using python for simplilearn-

Data Analysis (Western Michigan University)

In [55]:

```
# This Python 3 environment comes with many helpful analytics libraries installed
# It is defined by the kaggle/python Docker image: https://github.com/kaggle/docker-
# For example, here's several helpful packages to load

import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)

# Input data files are available in the read-only "../input/" directory
# For example, running this (by clicking run or pressing Shift+Enter) will list all

import os
for dirname, _, filenames in os.walk('/kaggle/input'):
    for filename in filenames:
        print(os.path.join(dirname, filename))

# You can write up to 20GB to the current directory (/kaggle/working/) that gets pre-
# You can also write temporary files to /kaggle/temp/, but they won't be saved outsi
```

/kaggle/input/house-loan-data-analysis/loan_data.csv

In [56]:

```
import pandas as pd
import sklearn
import numpy as np
import matplotlib.pyplot as plt
import os
import warnings
import seaborn as sns
from sklearn.preprocessing import OneHotEncoder
from sklearn.datasets import make_blobs
from sklearn.impute import SimpleImputer
from sklearn.pipeline import Pipeline
from sklearn.compose import ColumnTransformer
from sklearn.preprocessing import StandardScaler
from sklearn.svm import LinearSVC
from sklearn.metrics import roc_auc_score
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import roc_auc_score
from sklearn.calibration import CalibratedClassifierCV
from sklearn.metrics import confusion_matrix
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score
from sklearn.linear_model import SGDClassifier
import plotly.offline as py
import plotly.graph_objs as go
from plotly.offline import init_notebook_mode, iplot
from sklearn.model_selection import train_test_split
init_notebook_mode(connected=True)
import cufflinks as cf
cf.go_offline()
import pickle
import gc
import lightgbm as lgb
warnings.filterwarnings('ignore')
%matplotlib inline
```


In [57]:

```
house_loan=pd.read_csv('../input/house-loan-data-analysis/loan_data.csv')
house_loan.describe()
```

Out[57]:

	SK_ID_CURR	TARGET	CNT_CHILDREN	AMT_INCOME_TOTAL	AMT_CREDIT	AMT_ANNUITY
count	307511.000000	307511.000000	307511.000000	3.075110e+05	3.075110e+05	307499.0
mean	278180.518577	0.080729	0.417052	1.687979e+05	5.990260e+05	27108.5
std	102790.175348	0.272419	0.722121	2.371231e+05	4.024908e+05	14493.7
min	100002.000000	0.000000	0.000000	2.565000e+04	4.500000e+04	1615.5
25%	189145.500000	0.000000	0.000000	1.125000e+05	2.700000e+05	16524.0
50%	278202.000000	0.000000	0.000000	1.471500e+05	5.135310e+05	24903.0
75%	367142.500000	0.000000	1.000000	2.025000e+05	8.086500e+05	34596.0
max	456255.000000	1.000000	19.000000	1.170000e+08	4.050000e+06	258025.5

8 rows × 106 columns



In [58]:

house_loan.columns

```
Out[58]: Index(['SK_ID_CURR', 'TARGET', 'NAME_CONTRACT_TYPE', 'CODE_GENDER',
       'FLAG_OWN_CAR', 'FLAG_OWN_REALTY', 'CNT_CHILDREN', 'AMT_INCOME_TOTAL',
       'AMT_CREDIT', 'AMT_ANNUITY',
       ...
       'FLAG_DOCUMENT_18', 'FLAG_DOCUMENT_19', 'FLAG_DOCUMENT_20',
       'FLAG_DOCUMENT_21', 'AMT_REQ_CREDIT_BUREAU_HOUR',
       'AMT_REQ_CREDIT_BUREAU_DAY', 'AMT_REQ_CREDIT_BUREAU_WEEK',
       'AMT_REQ_CREDIT_BUREAU_MON', 'AMT_REQ_CREDIT_BUREAU_QRT',
       'AMT_REQ_CREDIT_BUREAU_YEAR'],
      dtype='object', length=122)
```

In [59]:

house_loan.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 307511 entries, 0 to 307510
Columns: 122 entries, SK_ID_CURR to AMT_REQ_CREDIT_BUREAU_YEAR
dtypes: float64(65), int64(41), object(16)
memory usage: 286.2+ MB
```

In [60]:

house_loan.isnull().sum()

```
Out[60]: SK_ID_CURR          0
TARGET            0
NAME_CONTRACT_TYPE  0
CODE_GENDER        0
FLAG_OWN_CAR        0
...
AMT_REQ_CREDIT_BUREAU_DAY  41519
AMT_REQ_CREDIT_BUREAU_WEEK 41519
AMT_REQ_CREDIT_BUREAU_MON   41519
AMT_REQ_CREDIT_BUREAU_QRT   41519
AMT_REQ_CREDIT_BUREAU_YEAR  41519
Length: 122, dtype: int64
```


In [61]:

house_loan.head()

	SK_ID_CURR	TARGET	NAME_CONTRACT_TYPE	CODE_GENDER	FLAG_OWN_CAR	FLAG_OWN_REALTY
0	100002	1	Cash loans	M	N	

SK_ID_CURR	TARGET	NAME_CONTRACT_TYPE	CODE_GENDER	FLAG_OWN_CAR	FLAG_OWN_REALM
1	100003	0	Cash loans	F	N
2	100004	0	Revolving loans	M	Y
3	100006	0	Cash loans	F	N
4	100007	0	Cash loans	M	N

5 rows × 122 columns



```
In [62]: defaulters=(house_loan.TARGET==1).sum()
payers=(house_loan.TARGET==0).sum()
print((defaulters/payers)*100)
```

8.781828601345662


```
In [63]: without_id=[column for column in house_loan.columns if column!='SK_ID_CURR']

#check for duplicate values
na=house_loan[house_loan.duplicated(subset=without_id,keep=False)]
print("Duplicates are: ",na.shape[0])
```

Duplicates are: 0

```
In [64]: house_loan.TARGET.value_counts().plot(kind='pie', autopct='%1.1f%%')
```


Out[64]: <AxesSubplot:ylabel='TARGET'>



```
In [65]: import matplotlib as plt
```

```
In [66]: shuffled_data=house_loan.sample(frac=1,random_state=3)
unpaid_home_loan=shuffled_data.loc[shuffled_data['TARGET']==1]
paid_home_loan=shuffled_data.loc[shuffled_data['TARGET']==0].sample(n=24825,random_state=3)
normalised_home_loan=pd.concat([unpaid_home_loan,paid_home_loan])
normalised_home_loan.TARGET.value_counts().plot(kind='pie', autopct="%1.1f%%")
```

Out[66]: <AxesSubplot:ylabel='TARGET'>



```
In [67]: import tensorflow as tf
```

```
In [68]: normalised_home_loan.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 49650 entries, 207339 to 121862
Columns: 122 entries, SK_ID_CURR to AMT_REQ_CREDIT_BUREAU_YEAR
dtypes: float64(65), int64(41), object(16)
memory usage: 46.6+ MB
```

```
In [69]: normalised_home_loan.head
```

```
Out[69]: <bound method NDFrame.head of
  SK_ID_CURR TARGET NAME_CONTRACT_TYPE CODE_GEN
  DER FLAG_OWN_CAR \
207339 340318 1 Cash loans F N
8756 110186 1 Cash loans M Y
230344 366811 1 Cash loans F N
178329 306645 1 Cash loans M Y
55586 164407 1 Cash loans M N
...
130947 251878 0 Cash loans F Y
40467 146875 0 Cash loans F N
187004 316791 0 Cash loans M N
131755 252811 0 Cash loans F N
121862 241287 0 Cash loans M N

  FLAG_OWN_REALTY CNT_CHILDREN AMT_INCOME_TOTAL AMT_CREDIT \
207339 N 0 112500.0 405000.0
8756 N 0 135000.0 544491.0
230344 Y 0 112500.0 225000.0
178329 Y 0 157500.0 595273.5
55586 N 0 157500.0 521451.0
...
130947 Y 0 135000.0 770913.0
40467 N 2 360000.0 260640.0
187004 Y 1 180000.0 688500.0
131755 Y 2 202500.0 312840.0
121862 N 0 58500.0 254700.0

  AMT_ANNUITY ... FLAG_DOCUMENT_18 FLAG_DOCUMENT_19 FLAG_DOCUMENT_20 \
207339 21969.0 ... 0 0 0
8756 17563.5 ... 0 0 0
230344 17905.5 ... 0 0 0
178329 29083.5 ... 0 0 0
55586 35406.0 ... 0 0 0
...
130947 24997.5 ... 0 0 0
40467 29475.0 ... 0 0 0
187004 22752.0 ... 0 0 0
```

```

131755      18090.0 ...      0      0      0
121862      13446.0 ...      0      0      0

      FLAG_DOCUMENT_21  AMT_REQ_CREDIT_BUREAU_HOUR  AMT_REQ_CREDIT_BUREAU_DAY \
207339          0                  0.0                0.0
8756          0                  0.0                0.0
230344          0                  NaN                NaN
178329          0                  NaN                NaN
55586          0                  0.0                0.0
...
130947          ...                ...                ...
40467          0                  0.0                0.0
187004          0                  0.0                0.0
131755          0                  0.0                0.0
121862          0                  0.0                0.0

      AMT_REQ_CREDIT_BUREAU_WEEK  AMT_REQ_CREDIT_BUREAU_MON \
207339          0.0                0.0
8756          0.0                0.0
230344          NaN                NaN
178329          NaN                NaN
55586          0.0                0.0
...
130947          ...                ...
40467          0.0                0.0
187004          0.0                0.0
131755          0.0                0.0
121862          0.0                0.0

      AMT_REQ_CREDIT_BUREAU_QRT  AMT_REQ_CREDIT_BUREAU_YEAR
207339          0.0                3.0
8756          0.0                0.0
230344          NaN                NaN
178329          NaN                NaN
55586          0.0                1.0
...
130947          ...                ...
40467          0.0                0.0
187004          0.0                0.0
131755          1.0                3.0
121862          0.0                0.0

```

[49650 rows x 122 columns]>

In [70]:

```
normalised_home_loan.dropna(axis=0)
normalised_home_loan.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 49650 entries, 207339 to 121862
Columns: 122 entries, SK_ID_CURR to AMT_REQ_CREDIT_BUREAU_YEAR
dtypes: float64(65), int64(41), object(16)
memory usage: 46.6+ MB
```

In [71]:

```
normalised_home_loan.isnull().sum()
```

Out[71]:

SK_ID_CURR	0
TARGET	0
NAME_CONTRACT_TYPE	0
CODE_GENDER	0
FLAG_OWN_CAR	0
...	
AMT_REQ_CREDIT_BUREAU_DAY	7648
AMT_REQ_CREDIT_BUREAU_WEEK	7648
AMT_REQ_CREDIT_BUREAU_MON	7648
AMT_REQ_CREDIT_BUREAU_QRT	7648
AMT_REQ_CREDIT_BUREAU_YEAR	7648

Length: 122, dtype: int64

This document is available free of charge on



In [72]: `#print(normalised_home_Loan.apply())`

In [73]: `print(pd.unique(normalised_home_loan.AMT_REQ_CREDIT_BUREAU_DAY))
print(pd.unique(normalised_home_loan.AMT_REQ_CREDIT_BUREAU_WEEK))
print(pd.unique(normalised_home_loan.AMT_REQ_CREDIT_BUREAU_MON))
print(pd.unique(normalised_home_loan.AMT_REQ_CREDIT_BUREAU_QRT))
print(pd.unique(normalised_home_loan.AMT_REQ_CREDIT_BUREAU_YEAR))`

```
[ 0. nan  1.  2.  4.  3.  9.]  
[ 0. nan  1.  2.  4.  3.  5.  6.]  
[ 0. nan  1.  3.  5.  9.  2.  6.  8.  4.  11. 12.  7. 13. 10. 17. 15. 14.  
 16. 18. 27.]  
[ 0. nan  2.  3.  1.  4.  5.  6. 19.  7.]  
[ 3.  0. nan  1.  5.  4.  2.  6.  7.  8.  9. 10. 14. 13. 12. 11. 22. 16.  
 23. 17.]
```

In [74]: `normalised_home_loan.dropna(axis=0)`

Out[74]:

	SK_ID_CURR	TARGET	NAME_CONTRACT_TYPE	CODE_GENDER	FLAG_OWN_CAR	FLAG_OWN
279124	423360	1	Cash loans	M	Y	
216116	350411	1	Cash loans	M	Y	
133687	255050	1	Cash loans	M	Y	
4159	104863	1	Cash loans	M	Y	
208602	341779	1	Cash loans	F	Y	
...
108677	226053	0	Cash loans	M	Y	
258603	399273	0	Revolving loans	M	Y	
51880	160079	0	Cash loans	M	Y	
282820	427561	0	Cash loans	F	Y	
207101	340051	0	Revolving loans	F	Y	

1230 rows × 122 columns


In [75]: `print(normalised_home_loan.info())
print(normalised_home_loan.isnull().sum())`

```
<class 'pandas.core.frame.DataFrame'>  
Int64Index: 49650 entries, 207339 to 121862  
Columns: 122 entries, SK_ID_CURR to AMT_REQ_CREDIT_BUREAU_YEAR  
dtypes: float64(65), int64(41), object(16)  
memory usage: 46.6+ MB  
None  
SK_ID_CURR                      0  
TARGET                          0  
NAME_CONTRACT_TYPE                0  
CODE_GENDER                       0  
FLAG_OWN_CAR                      0  
                                ...  
AMT_REQ_CREDIT_BUREAU_DAY        7648  
AMT_REQ_CREDIT_BUREAU_WEEK       7648  
AMT_REQ_CREDIT_BUREAU_MON        7648
```

```
AMT_REQ_CREDIT_BUREAU_QRT    7648  
AMT_REQ_CREDIT_BUREAU_YEAR   7648  
Length: 122, dtype: int64
```


```
In [76]: normalised_home_loan.TARGET.value_counts().plot(kind='pie', autopct="%1.1f%%")
```

```
Out[76]: <AxesSubplot:ylabel='TARGET'>
```




```
In [77]: normalised_home_loan.NAME_CONTRACT_TYPE.value_counts().plot(kind='pie', autopct="%1.1f%%")  
#high amount of cash Loans
```

```
Out[77]: <AxesSubplot:ylabel='NAME_CONTRACT_TYPE'>
```




```
In [78]: normalised_home_loan.CODE_GENDER.value_counts().plot(kind='pie', autopct="%1.1f%%")  
#roughly equal amount
```

```
Out[78]: <AxesSubplot:ylabel='CODE_GENDER'>
```




```
In [79]: normalised_home_loan.FLAG_own_car.value_counts().plot(kind='pie', autopct="%1.1f%%")
```

```
Out[79]: <AxesSubplot:ylabel='FLAG_own_CAR'>
```



```
In [80]: normalised_home_loan.CNT_CHILDREN.value_counts().plot(kind='pie', autopct="%1.1f%%")
```

```
Out[80]: <AxesSubplot:ylabel='CNT_CHILDREN'>
```




```
In [81]: #!pip install chart_studio
```

```
cf.set_config_file(theme='polar')
```

```
normalised_home_loan[normalised_home_loan['AMT_INCOME_TOTAL'] < 2000000]['AMT_INCOME_TOTAL']
xTitle = 'Total Income', yTitle = 'Count of applicants',
title='Distribution of AMT_INCOME_TOTAL')
```

Distribution of AMT_INCOME_TOTAL



```
In [82]: (normalised_home_loan[normalised_home_loan['AMT_INCOME_TOTAL']>1000000]['TARGET'].value
```

```
Out[82]: 0    64.864865
1    35.135135
Name: TARGET, dtype: float64
```

```
In [83]: #print((normalised_home_loan[normalised_home_loan['CNT_CHILDREN']>1]['TARGET'].value_
#      print((normalised_home_loan[normalised_home_loan['CNT_CHILDREN']>2]['TARGET'].value_
#      print((normalised_home_loan[normalised_home_loan['CNT_CHILDREN']>5]['TARGET'].value_
#      #as number of children is increasing lone defaulters are increasing
```

```
1    57.047872
0    42.952128
Name: TARGET, dtype: float64
1    81.818182
0    18.181818
Name: TARGET, dtype: float64
```

```
In [84]: print((normalised_home_loan[normalised_home_loan['FLAG_OWN_CAR']=='N']['TARGET'].value_
print((normalised_home_loan[normalised_home_loan['FLAG_OWN_CAR']=='Y']['TARGET'].value_
#people with own cars are slightly more likely to repay back the loan
```

```
1    51.350064
0    48.649936
Name: TARGET, dtype: float64
0    52.823962
1    47.176038
Name: TARGET, dtype: float64
```

```
In [85]: print((normalised_home_loan[normalised_home_loan['CODE_GENDER']=='M']['TARGET'].value_counts()))
print((normalised_home_loan[normalised_home_loan['CODE_GENDER']=='F']['TARGET'].value_counts()))

#men more likely to default in payment of loans
```

```
1    56.280372
0    43.719628
Name: TARGET, dtype: float64
0    53.867691
1    46.132309
Name: TARGET, dtype: float64
```

```
In [86]: print((normalised_home_loan[normalised_home_loan['NAME_CONTRACT_TYPE']=='Cash loans'].value_counts()))
print((normalised_home_loan[normalised_home_loan['NAME_CONTRACT_TYPE']=='Revolving loans'].value_counts()))

#cash Loans have a higher percent of defaulters
```

```
1    50.802923
0    49.197077
Name: TARGET, dtype: float64
0    59.309995
1    40.690005
Name: TARGET, dtype: float64
```

```
In [87]: normalised_home_loan=normalised_home_loan.sample(frac=1,random_state=5)
```

```
In [88]: from sklearn.preprocessing import OrdinalEncoder

ordenc=OrdinalEncoder()
normalised_home_loan['NAME_CONTRACT_TYPE_CODE']=ordenc.fit_transform(normalised_home_loan[['NAME_CONTRACT_TYPE']])
print(normalised_home_loan[['NAME_CONTRACT_TYPE','NAME_CONTRACT_TYPE_CODE']].head(20))
print(normalised_home_loan['NAME_CONTRACT_TYPE_CODE'].value_counts())
```

	NAME_CONTRACT_TYPE	NAME_CONTRACT_TYPE_CODE
302218	Cash loans	0.0
167526	Cash loans	0.0
159305	Cash loans	0.0
275427	Cash loans	0.0
8837	Cash loans	0.0
192094	Cash loans	0.0
235115	Revolving loans	1.0
79051	Cash loans	0.0
123267	Revolving loans	1.0
5517	Cash loans	0.0
128624	Cash loans	0.0
187583	Cash loans	0.0
143193	Cash loans	0.0
288269	Cash loans	0.0
44320	Cash loans	0.0
256898	Cash loans	0.0
118237	Cash loans	0.0
5980	Revolving loans	1.0
96475	Cash loans	0.0
249976	Cash loans	0.0
0.0	45708	
1.0	3942	

Name: NAME_CONTRACT_TYPE_CODE, dtype: int64

```
In [89]: normalised_home_loan['CODE_GENDER_CODE']=ordenc.fit_transform(normalised_home_loan[['CODE_GENDER']])
print(normalised_home_loan[['CODE_GENDER','CODE_GENDER_CODE']].head(20))
print(normalised_home_loan['CODE_GENDER_CODE'].value_counts())
```

```

CODE_GENDER  CODE_GENDER_CODE
302218      M          1.0
167526      F          0.0
159305      M          1.0
275427      F          0.0
8837        M          1.0
192094      M          1.0
235115      F          0.0
79051       F          0.0
123267      M          1.0
5517        F          0.0
128624      M          1.0
187583      F          0.0
143193      M          1.0
288269      F          0.0
44320       F          0.0
256898      F          0.0
118237      F          0.0
5980        M          1.0
96475       F          0.0
249976      F          0.0
0.0         30716
1.0         18932
2.0         2
Name: CODE_GENDER_CODE, dtype: int64

```

In [90]:

```
#2 other values in code_gender
normalised_home_loan.loc[normalised_home_loan['CODE_GENDER_CODE'] == 2]
```

Out[90]:

	SK_ID_CURR	TARGET	NAME_CONTRACT_TYPE	CODE_GENDER	FLAG_OWN_CAR	FLAG_OWN
	83382	196708	0	Revolving loans	XNA	N
	189640	319880	0	Revolving loans	XNA	Y

2 rows × 124 columns



In [91]:

```
normalised_home_loan['FLAG_OWN_CAR_CODE']=ordenc.fit_transform(normalised_home_loan[
print(normalised_home_loan[['FLAG_OWN_CAR','FLAG_OWN_CAR_CODE']].head(20))
print(normalised_home_loan['FLAG_OWN_CAR_CODE'].value_counts())
```

FLAG_OWN_CAR	FLAG_OWN_CAR_CODE	
302218	N	0.0
167526	N	0.0
159305	N	0.0
275427	N	0.0
8837	N	0.0
192094	N	0.0
235115	N	0.0
79051	N	0.0
123267	N	0.0
5517	N	0.0
128624	N	0.0
187583	N	0.0
143193	N	0.0
288269	Y	1.0
44320	Y	1.0
256898	N	0.0
118237	N	0.0
5980	Y	1.0
96475	N	0.0
249976	N	0.0
0.0	33591	

```
1.0    16059
Name: FLAG_OWN_CAR_CODE, dtype: int64
```

```
In [92]: normalised_home_loan['CNT_CHILDREN_CODE']=ordenc.fit_transform(normalised_home_loan[['CNT_CHILDREN_CODE','CNT_CHILDREN']])
print(normalised_home_loan[['CNT_CHILDREN_CODE','CNT_CHILDREN']].head(20))
print(normalised_home_loan['CNT_CHILDREN_CODE'].value_counts())
```

	CNT_CHILDREN_CODE	CNT_CHILDREN
302218	0.0	0
167526	0.0	0
159305	2.0	2
275427	0.0	0
8837	0.0	0
192094	0.0	0
235115	0.0	0
79051	0.0	0
123267	1.0	1
5517	0.0	0
128624	0.0	0
187583	1.0	1
143193	0.0	0
288269	0.0	0
44320	0.0	0
256898	0.0	0
118237	2.0	2
5980	0.0	0
96475	0.0	0
249976	0.0	0
0.0	34073	
1.0	10381	
2.0	4444	
3.0	642	
4.0	89	
5.0	10	
6.0	6	
8.0	2	
7.0	1	
9.0	1	
10.0	1	

```
Name: CNT_CHILDREN_CODE, dtype: int64
```

```
In [93]: normalised_home_loan=normalised_home_loan.sample(frac=1,random_state=45)
```

```
In [94]: normalised_home_loan['TARGET'].value_counts()
```

```
Out[94]: 0    24825
1    24825
Name: TARGET, dtype: int64
```

```
In [95]: y=normalised_home_loan.TARGET
```

```
In [96]: #y=y.sample(frac=1,random_state=45)
```

```
In [97]: normalised_home_loan_features=['SK_ID_CURR','NAME_CONTRACT_TYPE_CODE','CNT_CHILDREN_
```

```
In [98]: from sklearn.model_selection import train_test_split
```

```
In [99]: X=normalised_home_loan[normalised_home_loan_features]
```

```
In [100...]: #X=X.sample(frac=1,random_state=45)
```

```
In [101...]: blobs_random_seed = 42
centers = [(0,0), (5,5)]
cluster_std = 1
frac_test_split = 0.33
num_features_for_samples = 2
num_samples_total = 49650


# Generate data
inputs, targets = make_blobs(n_samples = num_samples_total, centers = centers, n_features = num_features_for_samples, cluster_std = cluster_std, random_state = blobs_random_seed)

X_train,X_test,y_train,y_test=train_test_split(inputs,targets,test_size=0.33,random_state=45)
```

```
In [102...]: print(X_train.shape, X_test.shape, y_train.shape, y_test.shape)
```

(33265, 2) (16385, 2) (33265,) (16385,)

```
In [103...]: plt.pyplot.scatter(X_train[:,0], X_train[:,1])
plt.pyplot.title('Linearly separable data')
plt.pyplot.xlabel('X1')
plt.pyplot.ylabel('X2')
plt.pyplot.show()
```



```
In [104...]: from sklearn import svm
from sklearn.metrics import plot_confusion_matrix
```

```
In [105...]: clf=SVM(kernel='linear')
```


```
In [106...]: clf=clf.fit(X_train,y_train)
```

```
In [107...]: predictions = clf.predict(X_test)
```

```
# Generate confusion matrix
matrix = plot_confusion_matrix(clf, X_test, y_test,
                               cmap=plt.cm.Blues)
```

This document is available free of charge on

```
normalize='true')
plt.pyplot.title('Confusion matrix for our classifier')
plt.pyplot.show(matrix)
plt.pyplot.show()
```



In [118]:

```
from sklearn.metrics import precision_score, recall_score, f1_score
```

In [119]:


```
print(precision_score(y_test, predictions))
print(recall_score(y_test, predictions))
print(f1_score(y_test, predictions, average=None))
```

```
0.9995034140285537
0.9998758072528564
[0.99969993 0.99968958]
```

In [108]:


```
support_vectors = clf.support_vectors_

# Visualize support vectors
plt.pyplot.scatter(X_train[:,0], X_train[:,1])
plt.pyplot.scatter(support_vectors[:,0], support_vectors[:,1], color='red')
plt.pyplot.title('Linearly separable data with support vectors')
plt.pyplot.xlabel('X1')
plt.pyplot.ylabel('X2')
plt.pyplot.show()
```



```
In [109...]: from mlxtend.plotting import plot_decision_regions
```

```
In [112...]: plot_decision_regions(X_test, y_test, clf=clf, legend=2)
plt.pyplot.show()
```



```
In [ ]:
```